

NERSC User's Group (NUG) Community Call:

Using CUDA with Python on
Perlmutter@NERSC



15 August, 2024

Today's Pipeline

- Front Matter/Getting Started
- Jupyter Primer
- Interactive Jupyter Notebook
 - Python Review
 - Cuda with Python
- Introduction to CuPy
- Advanced Features of CuPy

Some Logistics

- In-person attendees please also join Zoom for full participation
- Please change your name in Zoom session
 - to: first_name last_name
 - Click “Participants”, then “More” next to your name to rename
- Click the CC button to toggle captions and View Full Transcript
- Session is being recorded
- Users are muted upon joining Zoom
 - Feel free to unmute and ask questions or ask in GDoc below
- GDoc is used for Q&A (instead of Zoom chat)
- Please answer a short survey afterward

Some Logistics

- Slides and videos will be available on Grads@NERSC page and NERSC Training Event page
- Introduction to CUDA Programming Training
 - <https://www.nersc.gov/users/training/training-materials/>
 - previous training materials available
 - 13-Part Detailed CUDA Training

Hands-on Exercises on Perlmutter

ssh <user>@perlmutter.nersc.gov

- **% cd \$SCRATCH**
- **% git clone** https://github.com/charlesWLivelyPhD/NERSC_CUDAPYTHON_August2024
 - Downloads all exercises (and answers!)
- References
 - <https://docs.nersc.gov/jobs/>
 - <https://docs.nersc.gov/jobs/examples/#interactive>

Using Perlmutter Compute Node Reservations

- Existing NERSC users (at time of registration) have been added to “`ntrain3`” project
- Non-NERSC users have received email instructions on apply for a training account
 - Please let us know if you need one
- Perlmutter node reservations: 10:30 am - 4:30 pm PDT today
 - `--reservation=august_nug -A ntrain3 -C gpu`
(add `-q shared -c 32 -G 1` for shared)
for sbatch or salloc sessions
 - No need to use `--reservation` or `-A` when outside of the reservation hours



NERSC Code of Conduct

Team Science

Service

Trust

Innovation

Respect




■ We agree to **work together professionally and productively** towards our shared goals while respecting each other's differences and ideas.

■ We should all feel free to speak up to maintain this environment and remember there are resources available to **report violations** to foster an inclusive, collaborative environment.
Email nersc-training@lbl.gov for any concerns

<https://www.nersc.gov/nersc-code-of-conduct>





Using Python on GPUs and Jupyter on Perlmutter@NERSC



Getting started with GPUs in Python

- NumPy and SciPy do not utilize GPUs out of the box
- There are many Python GPU frameworks out there:
 - “drop in” replacements for numpy, scipy, pandas, scikit-learn, etc
 - **CuPy**, **RAPIDS**
 - “machine learning” libraries that also support general GPU computing
 - **PyTorch**, **TensorFlow**, **JAX**
 - “I want to write my own GPU kernels”
 - **Numba**, **CUDA Python**
 - multi-gpu / multi-node / distributed memory:
 - **mpi4py+X**, **dask**, **cuNumeric**
- Many of these GPU libraries have adopted the CUDAArray Interface which makes it easier to pass array-like objects stored in GPU memory between the libraries
- There is also effort in the community to standardize around a common Python array API



```
numpy:      mean(a, axis=None, dtype=None, out=None, keepdims=<no value>)
dask.array: mean(a, axis=None, dtype=None, out=None, keepdims=<no value>)
cupy:      mean(a, axis=None, dtype=None, out=None, keepdims=False)
jax.numpy: mean(a, axis=None, dtype=None, out=None, keepdims=False)
mxnet.np:  mean(a, axis=None, dtype=None, out=None, keepdims=False)
sparse:    s.mean(axis=None, keepdims=False, dtype=None, out=None)
torch:     mean(input, dim, keepdim=False, out=None)
tensorflow: reduce_mean(input_tensor, axis=None, keepdims=None, name=None,
                        reduction_indices=None, keep_dims=None)
```

cuda toolkit dependency via module

```
> module load conda
```

Note: cuda toolkit module is loaded by default
Current default version is cuda toolkit/12.2

```
> conda create --name cupy-demo python=3.12 numpy scipy
```

```
> conda activate cupy-demo
```

```
> pip install cupy-cuda11X
```

Check your package documentation to see
cuda toolkit compatibility requirements

```
> python
```

```
>>> import cupy as cp
```

```
>>> print(cp.array([1, 2, 3]))
```

```
[1 2 3]
```

<https://docs.nersc.gov/development/languages/python/using-python-perlmutter/>

cuda toolkit dependency via conda

```
> module load conda
> module unload cudatoolkit
> conda create --name cupy-demo python=3.11 numpy scipy
> conda activate cupy-demo
> conda install -c conda-forge cupy
> python
>>> import cupy as cp
>>> print(cp.array([1, 2, 3]))
[1 2 3]
```

cupy conda-forge package will pull cudatoolkit dependencies into conda env

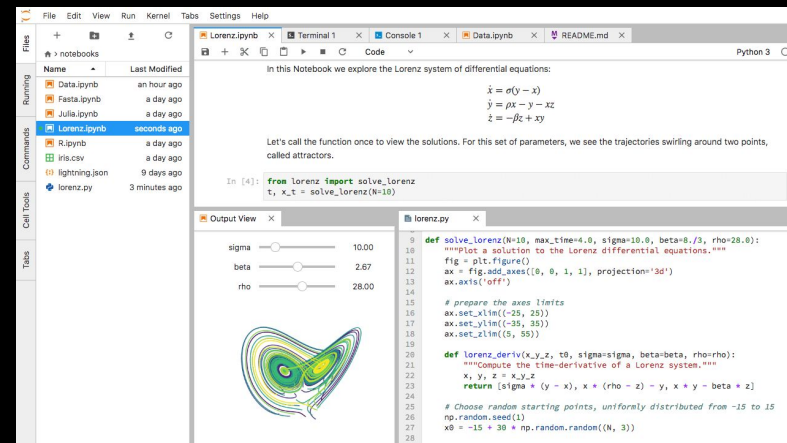
cupy conda-forge package will pull cudatoolkit dependencies into conda env

<https://docs.nersc.gov/development/languages/python/using-python-perlmutter/>

What is Jupyter?

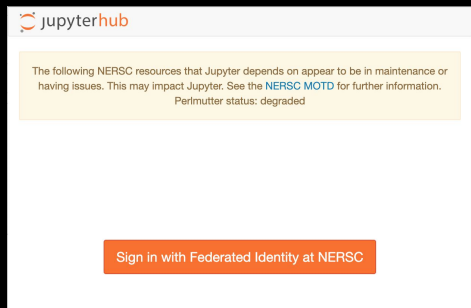


- At NERSC, we say “Jupyter” in reference to a collection of many things
 - Access shareable Jupyter “notebooks” via JupyterHub
- What can I put in a Jupyter notebook?
 - Live code
 - Equations
 - Visualizations
 - Narrative text
 - Interactive widgets
- What applications would I use a notebook for?
 - Data cleaning and data transformation
 - Numerical simulation
 - Statistical modeling
 - Data visualization
 - Machine learning
 - Workflows and analytics frameworks




How Do I Use Jupyter at NERSC?

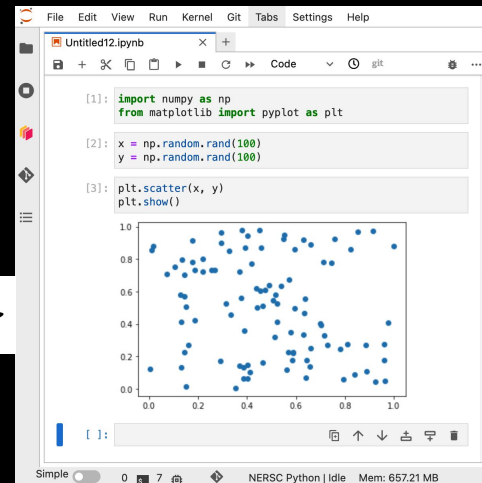
- <https://jupyter.nersc.gov>



Authenticate



Choose



Go!

How Do I Choose a Notebook Server to Spawn?

Shared CPU:

Exclusive CPU/GPU:

Configurable Job:

`!sbatch`

	Shared CPU Node	Exclusive CPU Node	Exclusive GPU Node	Configurable Job
Perlmutter	start	start	start	start
<i>Resources</i>	Use a node shared with other users' notebooks but outside the batch queues.	Use your own node within a job allocation using defaults.		Use multiple compute nodes with specialized settings.
<i>Use Cases</i>	Visualization and analytics that are not memory intensive and can run on just a few cores.	Visualization, analytics, machine learning that is compute or memory intensive but can be done on a single node.		Multi-node analytics jobs, jobs in reservations, custom project charging, and more.

**Shared = other users
and processes on the
same node**

**Exclusive and configurable =
compute nodes just for your
notebook and processes**

Configurable Job Settings


[Logout](#)

Configurable Job

start

Use multiple compute nodes with specialized settings.

Multi-node analytics jobs, jobs in reservations, custom project charging, and more.



Server Options

Account ("_g" suffix will be added as needed):

nstaff

Constraint:

gpu

QOS:

jupyter

cpus-per-task (node has 128 cpus):

128

gpus-per-task (node has 4 GPUs):

4

nodes (maximum of 4 for jupyter QOS):

1

ntasks-per-node:

1


Reservation:

(None)

time (time limit in minutes):


360

Start



Launch Interactive Jupyter Notebook





Introduction to CuPY



CuPy is...

a library to provide NumPy-compatible features with GPU



```
import numpy as np
X_cpu = np.zeros((10,))
W_cpu = np.zeros((10, 5))
y_cpu = np.dot(x_cpu,
W_cpu)
```

y_cpu = cp.asnumpy(y_gpu)




```
import cupy as cp
x_gpu = cp.zeros((10,))
W_gpu = cp.zeros((10, 5))
y_gpu = cp.dot(x_gpu,
W_gpu)
```

y_gpu = cp.asarray(y_cpu)

Comparing NumPy and CuPy

```
import numpy as np
```



NumPy

```
X_cpu =  
np.zeros((10,))  
W_cpu = np.zeros((10, 5))  
y_cpu = np.dot(x_cpu, W_cpu)
```

```
import cupy as cp
```



CuPy

```
x_gpu = cp.zeros((10,))  
W_gpu = cp.zeros((10, 5))  
y_gpu = cp.dot(x_gpu,  
W_gpu)
```



```
for xp in [np, cp]:  
    x = xp.zeros((10,))  
    W = xp.zeros((10, 5))  
    y = xp.dot(x, W)
```

Support both CPU and GPU with the same code!

Why develop CuPy? (1)

- Chainer functions had separate implementations in NumPy and PyCUDA to support both CPU and GPU

Even writing simple functions like “Add” or “Concat” took several lines...

```
7  _args = 'const float* x, float* y, int cdimx, int cdimy, int rdim, int coffset'
8  _preamble = '''
9  #define COPY(statement) \
10      int l    = i / (rdim * cdimx); \
11      int c    = i / rdim % cdimx + coffset; \
12      int r    = i % rdim; \
13      int idx = r + rdim * (c + cdimy * l); \
14      statement;
15  '''
16
17
18  class Concat(function.Function):
19
20      """Concatenate multiple tensors towards specified axis."""
21
```

Why develop CuPy? (2)

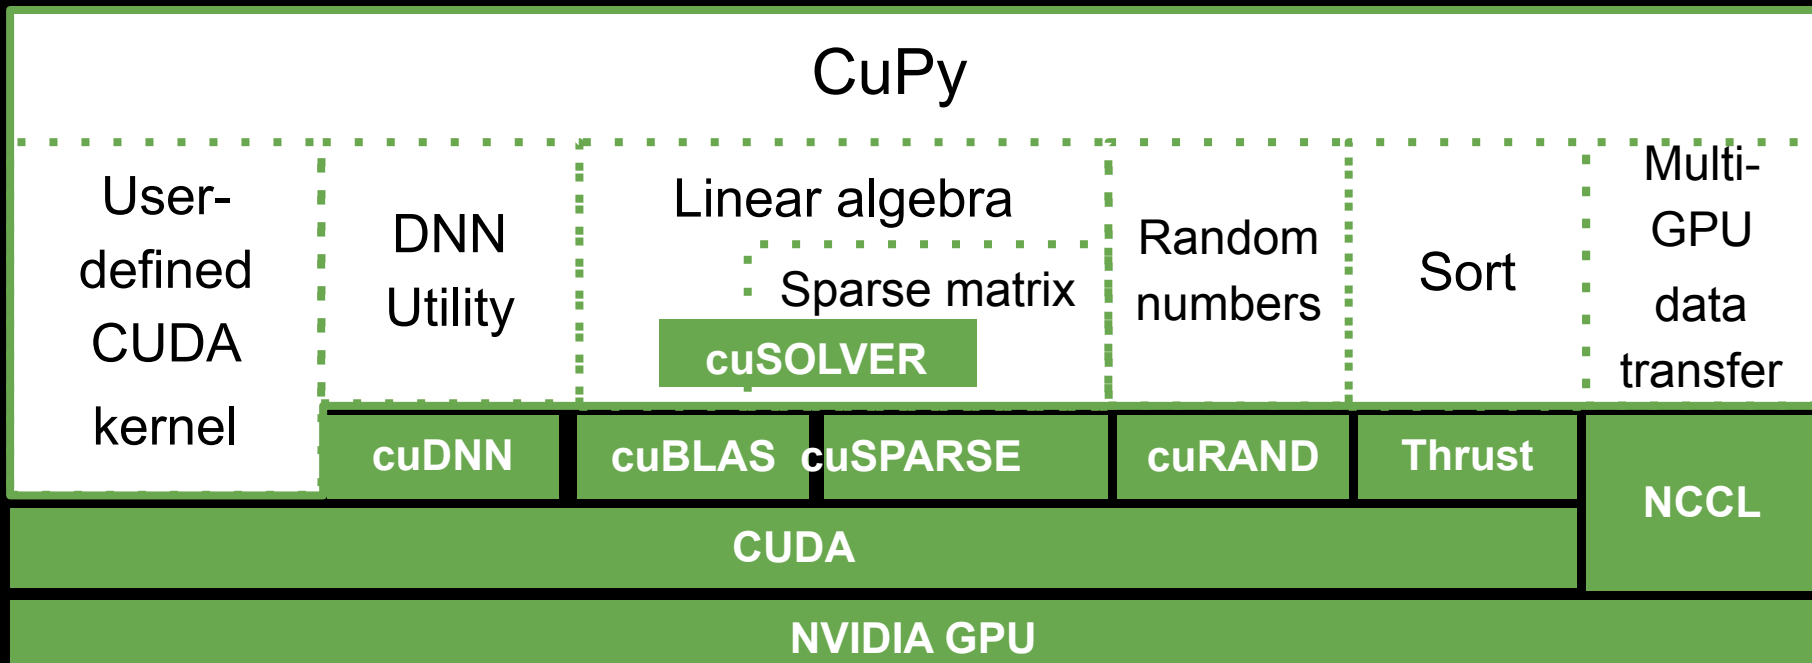
- Needed a NumPy-compatible GPU array library
 - NumPy is complicated
 - dtypes
 - Broadcast
 - Indexing

Why develop CuPy? (3)

- There was no convenient library
 - gnumpy
 - Consists of a single file which has 1000 lines of code
 - Not currently maintained
 - CUDA-based NumPy
 - No pip package is provided

⇒NVIDIA Needed to develop it

Inside CuPy



NumPy compatible features

Data types (dtypes)

bool_, int8, int16, int32, int64, uint8,
uint16,
uint32, uint64, float16, float32, float64,
complex64, and complex128

All basic indexing

indexing by ints, slices, newaxes, and
Ellipsis

Most of advanced indexing

except indexing patterns with boolean
masks

Most of the array creation routines

empty, ones_like, diag, etc...

Most of the array manipulation routines

reshape, rollaxis, concatenate, etc...

All operators with broadcasting

All universal functions for element-wise operations

except those for complex numbers

Linear algebra functions accelerated by cuBLAS

including product: dot, matmul, etc...

including decomposition: cholesky,
svd, etc...

Reduction along axes

sum, max, argmax, etc...


Sort operations implemented by Thrust

sort, argsort, and lexsort

Sparse matrix accelerated by cuSPARSE

New features after CuPy v2

- Narrowed the gap with NumPy
- Speedup: Cythonized, Improved MemoryPool
- CUDA Stream support
- Added supported functions
 - From NumPy
 - Sparse Matrix, FFT, scipy ndimage support



Advanced Features in CuPY



Overview of Advanced Features

- Kernel Fusion
- Unified Memory
- Custom Kernels
- Compatibility with other libraries
 - SciPy-compatible features
 - Direct use of NumPy functions via `array_interface`
 - Numba
 - PyTorch via DLPack
 - cuDF / cuML

Fusion: fuse kernels for further speedup!

```
a = numpy.float32(2.0)
x = xp.ones((1024, size), 'f')
y = xp.ones((1024, size), 'f')
```

```
def saxpy(a, x, y):
    return a * x + y
```

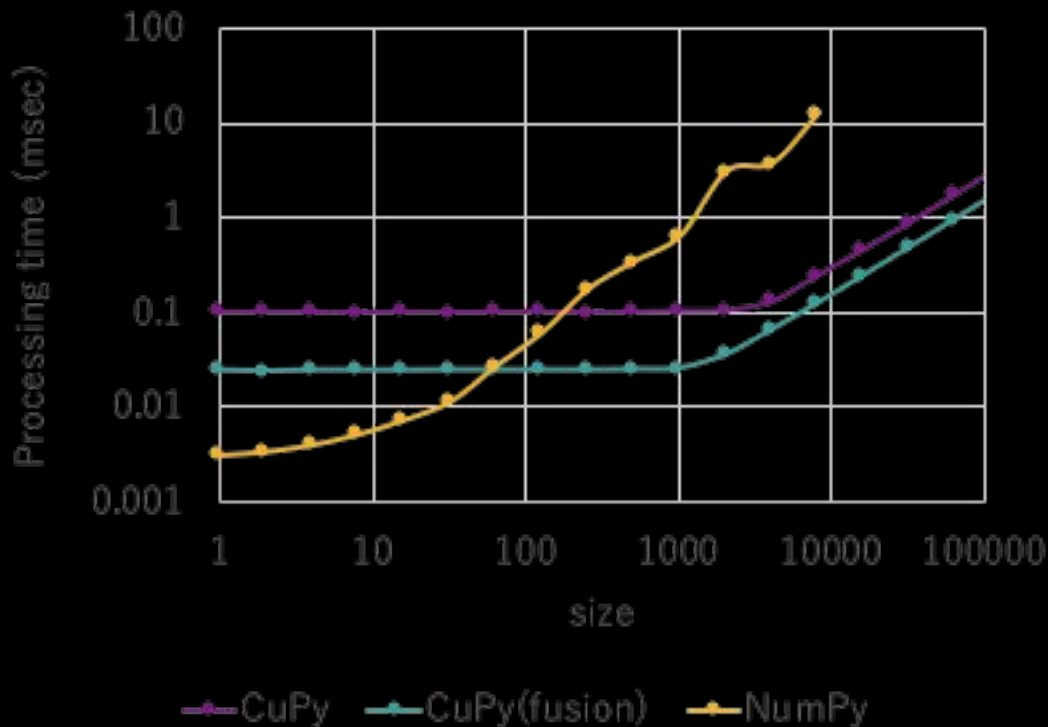
```
saxpy(a, x, y) # target
```

↓

```
@cupy.fuse()
```

```
def saxpy(a, x, y):
    return a * x + y
```

```
saxpy(a, x, y) # target
```



- Speedup function calls
- Reduce memory consumption
- Relax the bandwidth bottleneck

Limitations of `@cupy.fuse()`

- Only element-wise and reduction operations are supported
- Other operations like `cupy.matmul()` and `cupy.reshape()` are not yet supported

You want to save GPU memory?

```
import cupy as cp
size = 32768
a = cp.ones((size, size)) # 8GB
b = cp.ones((size, size)) # 8GB
cp.dot(a, b) # 8GB
```



Traceback (most recent call last):

...

cupy.cuda.memory.OutOfMemoryError: out of memory
to allocate 8589934592 bytes (total 17179869184
bytes)

Try Unified Memory! (Supported only on V100)

- Just edit 2 lines to enable **unified memory**

```
import cupy as cp

pool =
cp.cuda.MemoryPool(cp.cuda.malloc_managed)
cp.cuda.set_allocator(pool.malloc)
size = 32768
a = cp.ones((size, size)) # 8GB
b = cp.ones((size, size)) # 8GB
cp.dot(a, b)               # 8GB
```

- CuPy provides classes to compile **your own CUDA kernel**:
 - ElementwiseKernel
 - ReductionKernel
 - **RawKernel(from v5)**
 - For CUDA experts who love to write everything by themselves
 - Compiled with NVRTC

Basic usage of ElementwiseKernel

```
squared_diff = cp.ElementwiseKernel(  
    'float32 x, float32 y',    # input params  
    'float32 z',              # output params  
    'z = (x - y) * (x - y)',   # element-wise operation  
    'squared_diff'            # the name of this kernel  
)  
  
x = cp.arange(10, dtype=np.float32).reshape(2, 5)  
y = cp.arange(5, dtype=np.float32)  
  
squared_diff(x, y)
```

Type-generic kernels

```
squared_diff_generic =  
cp.ElementwiseKernel(  
    'T x, T y',                                     # input params  
    'T z',                                           # output params  
    'z = (x - y) * (x - y)',                         # element-wise operation  
    'squared_diff'                                  # the name of this  
                                                    kernel  
)  
  
x = cp.arange(10, dtype=np.float32).reshape(2,  
5) y = cp.arange(5, dtype=np.float32)  
  
squared_diff_generic(x, y)
```

Type-generic kernels

```
squared_diff_generic = cp.ElementwiseKernel(
    'T x, T y',
    'T z',
    '''
        T diff = x - y; z =
        diff * diff;
    ''',
    'squared_diff_generic')

x = cp.arange(10, dtype=np.float32).reshape(2, 5)
y = cp.arange(5, dtype=np.float32)

squared_diff_generic(x, y)
```

Manual indexing with `raw`

```
add_reverse = cp.ElementwiseKernel(  
    'T x, raw T y',           # input params  
    'T z',                     # output params  
    'z = x + y[_ind.size() - i - 1]', # element-wise operation  
    'add_reverse'              # the name of this kernel  
)
```

```
x = cp.arange(5,  
dtype=np.float32) y =  
cp.arange(5, dtype=np.float32)
```

```
add_reverse(x, y)
```


Reduction Kernel

```
l2norm_kernel = cp.ReductionKernel(  
    'T x',          # input array  
    'T y',          # output array  
    'x * x',        # map  
    'a + b',        # reduce  
    'y = sqrt(a)',  # post-reduction map  
    '0',            # identity value  
    'l2norm'        # kernel name  
)  
x = cp.arange(1000, dtype=np.float32).reshape(20,  
50) l2norm_kernel(x, axis=1)
```

=> This is same as : `cp.sqrt((x * x).sum(axis=1))`
but much faster!

How a RawKernel looks...

```
import cupy as cp

square_kernel = cp.RawKernel(r'''
extern "C" __global void my_square(long long* x) {
    int tid = threadIdx.x;
    x[tid] *= x[tid];
}
''', name='my_square')

x = cp.arange(5)
square_kernel(grid=(1,), block=(5,), args=(x,))
print(x)    # [ 0  1  4  9 16]
```

Now onto the Interactive Notebooks

